

Classical machine learning models

Amira Abbas

IBM Quantum, University of KwaZulu-Natal



Recap

DATA

Recap

DATA

MODEL

$$f(X_{train}; \vec{\theta})$$

Recap

DATA

A 10x10 grid of handwritten digits from 0 to 9, arranged in two rows. The first row contains digits 0 through 4, and the second row contains digits 5 through 9. Each digit is written in a different style and orientation.

MODEL

$$f(X_{train}; \vec{\theta})$$

COST

$$C(f(X_{train}; \vec{\theta}), y)$$

Recap

DATA

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3
4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4
5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5
6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6
7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7
8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8
9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9

MODEL

$$f(X_{train}; \vec{\theta})$$

↓

↑

COST

$$C(f(X_{train}; \vec{\theta}), y)$$

OPTIMIZATION

$$\frac{\partial C(f(X_{train}; \vec{\theta}), y)}{\partial \vec{\theta}}$$

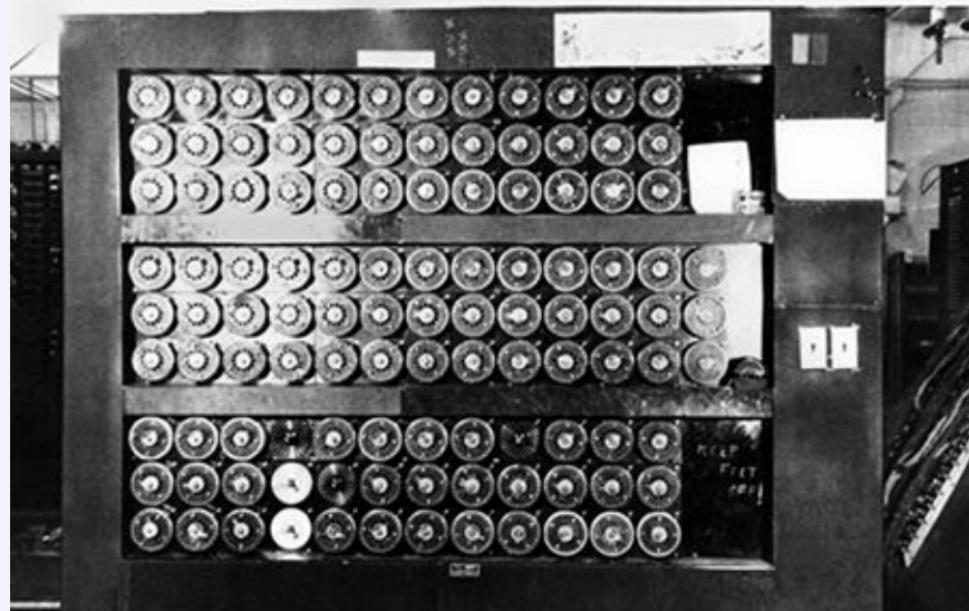
Recap

MODEL

$$f(X_{train}; \vec{\theta})$$

History of machine learning

- 1950: Turing's Learning Machine



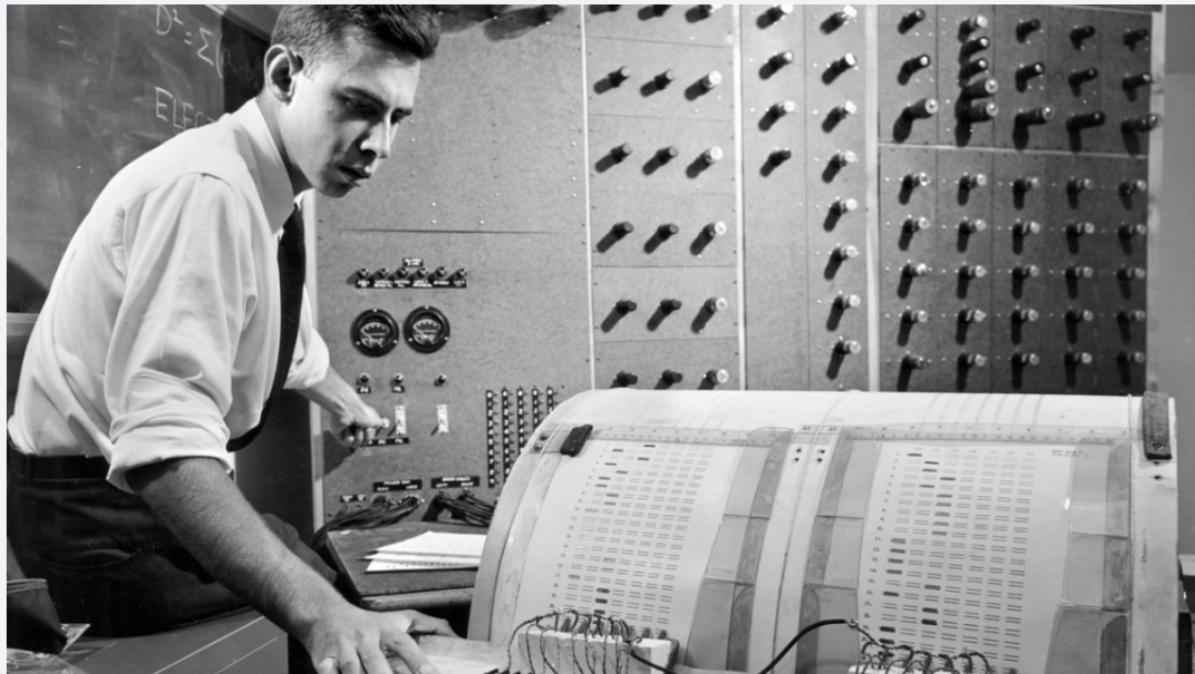
History of machine learning

- Machine learning was coined in 1952 by Arthur Samuel, an American IBMer



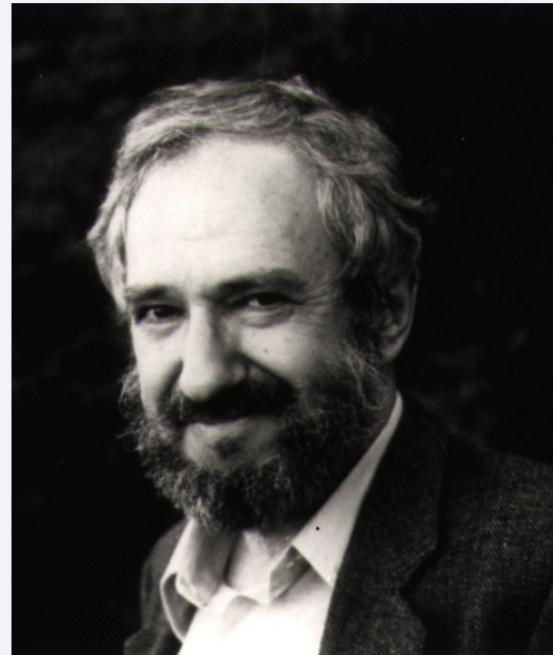
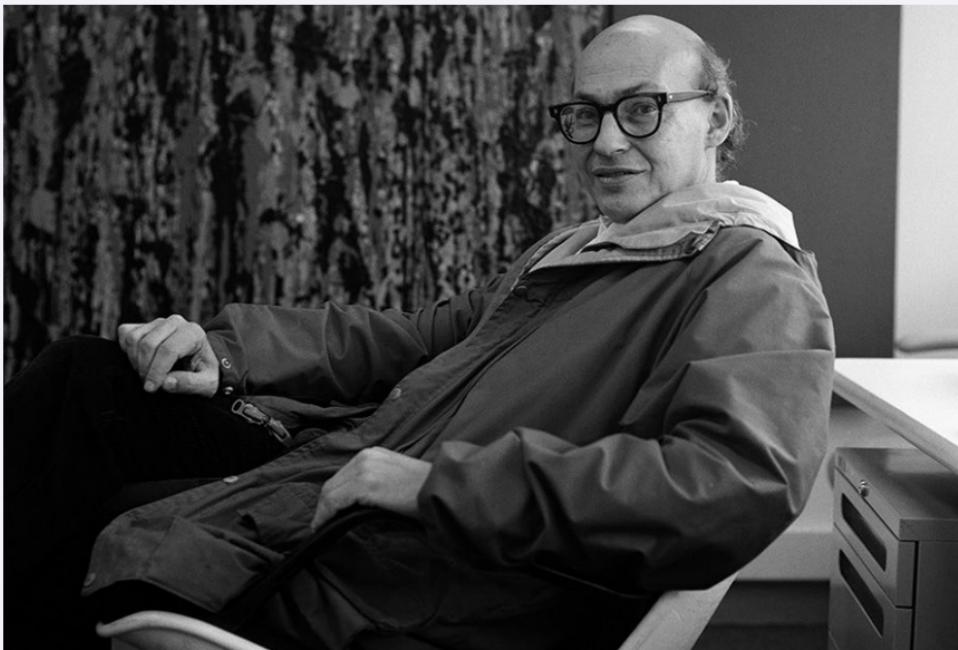
History of machine learning

- 1957: Enter the perceptron model created by Frank Rosenblatt, Cornell



History of machine learning

- 1969: Limitations of neural networks, Marvin Minsky & Seymour Papert



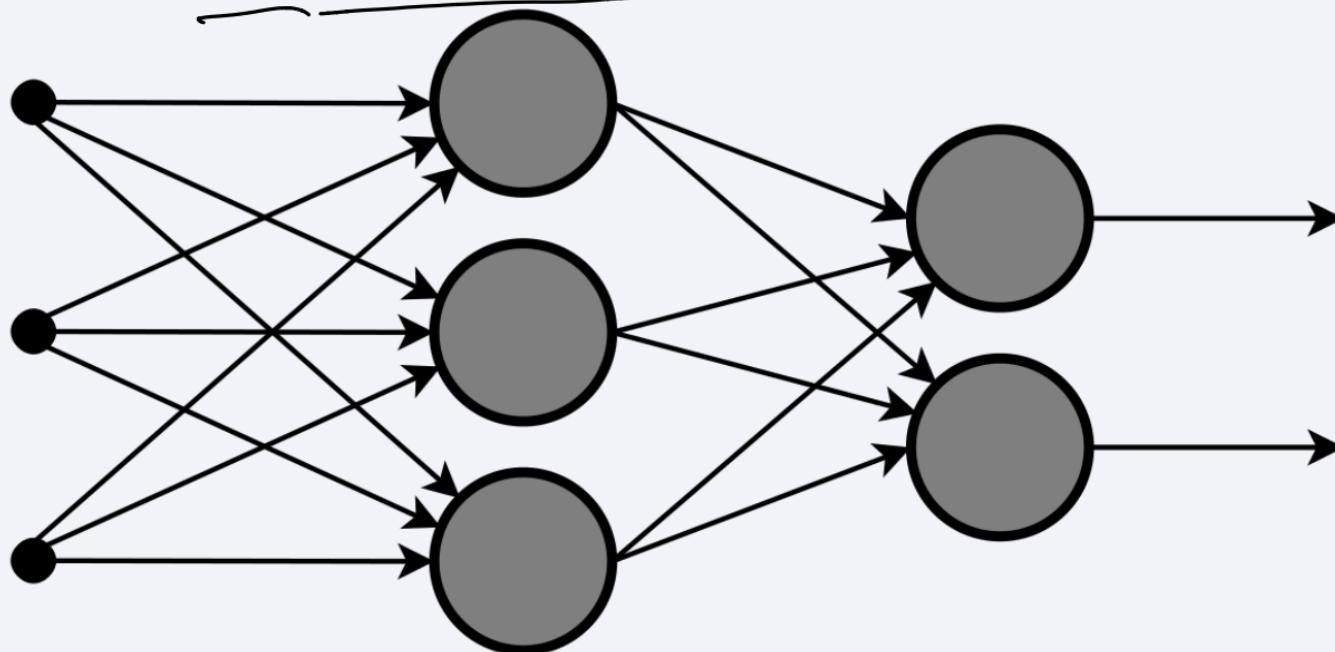
History of machine learning

- 1970's: AI winter



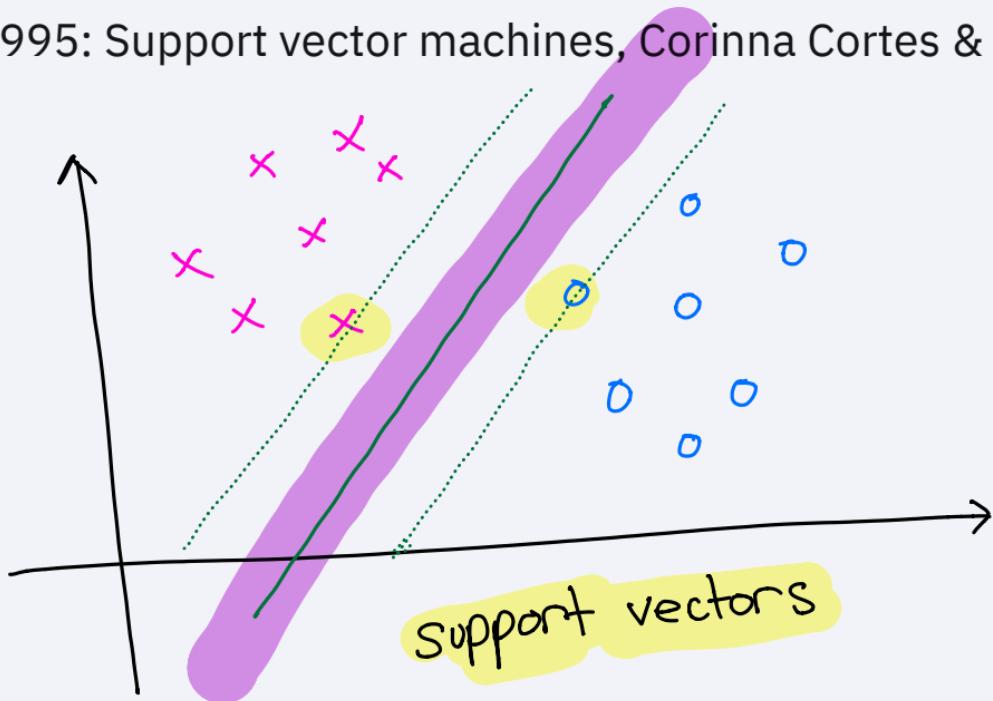
History of machine learning

- 1970 & 1986: Backpropagation, Seppo Linnainmaa & Geoffrey Hinton



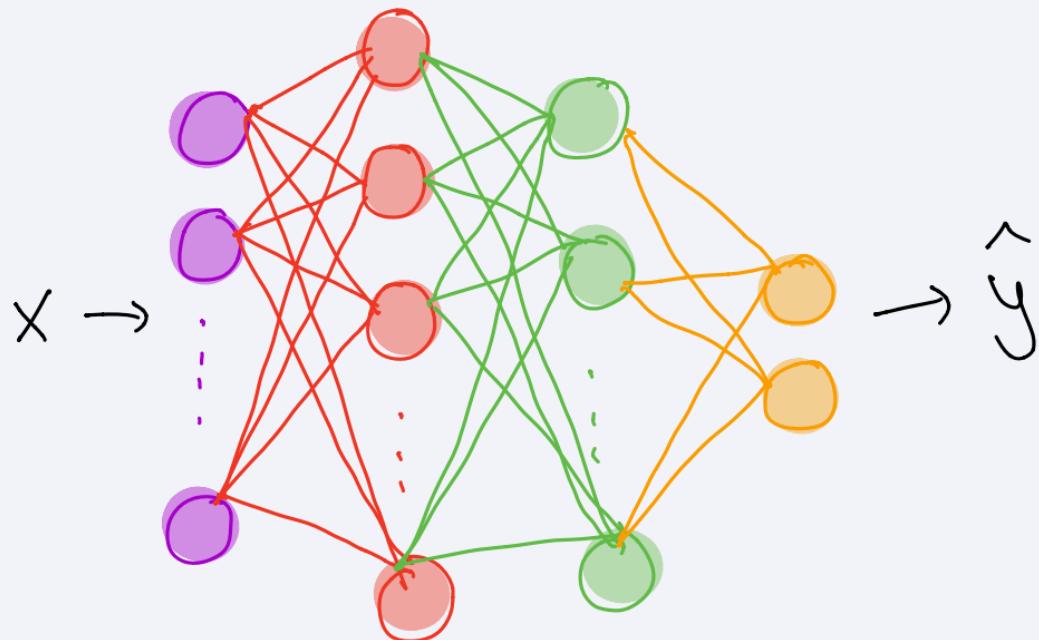
History of machine learning

- 1995: Support vector machines, Corinna Cortes & Vladimir Vapnik



History of machine learning

- 2010s: Deep learning became feasible in practice



Linear models

Linear models

Linear models

1-dimensional data, regression problem

$$\vec{x}_i = [-] \quad \text{dataset} = (\vec{x}_1, \dots, \vec{x}_n)$$

d-dimensional data

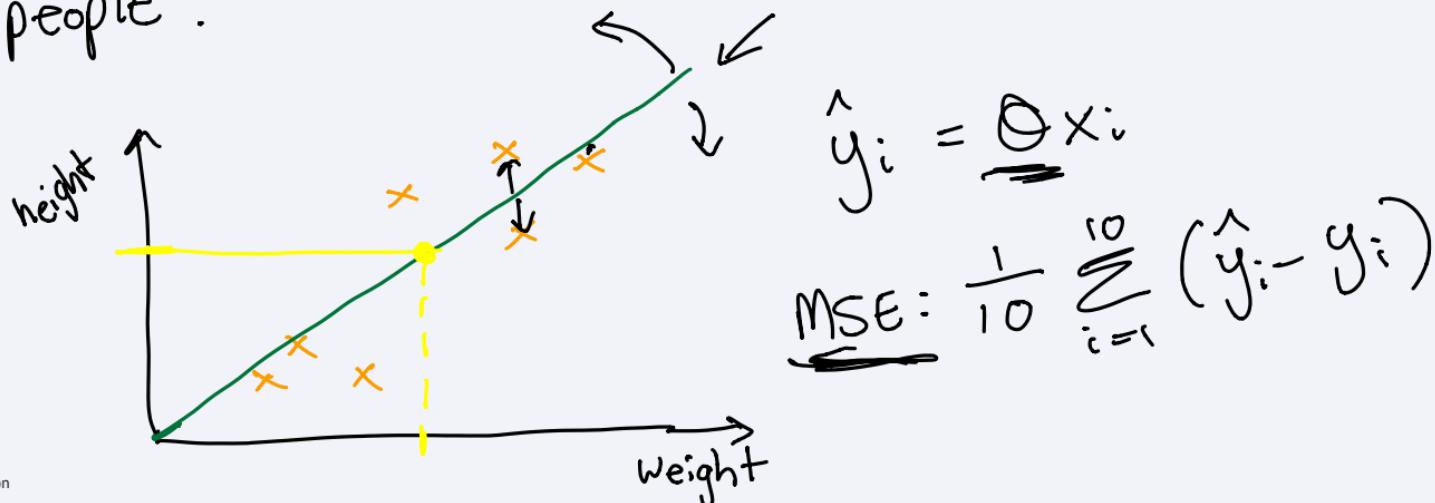
$$\vec{x}_i = \begin{bmatrix} 1 \\ \vdots \\ d \end{bmatrix} \quad d \rightarrow \text{features}$$

Linear models

1-dimensional data, regression problem

Task: predict the height of a person, based on their weight

$n=10$ samples , we know heights $\overline{(y_i)}$ and weights (x_i) of 10 people .



Linear models

Higher dimensional data, regression problem

Task: predict the height of a person, based on their weight, gender, height of their parents, etc.

$$\vec{x} = \begin{bmatrix} \text{weight} & x_1 \\ \text{gender} & x_2 \\ \text{height parent 1} & x_3 \end{bmatrix}$$

$$\hat{y} = \vec{\theta}^T \vec{x}$$

[$\theta_0, \theta_1, \theta_2, \theta_3$]

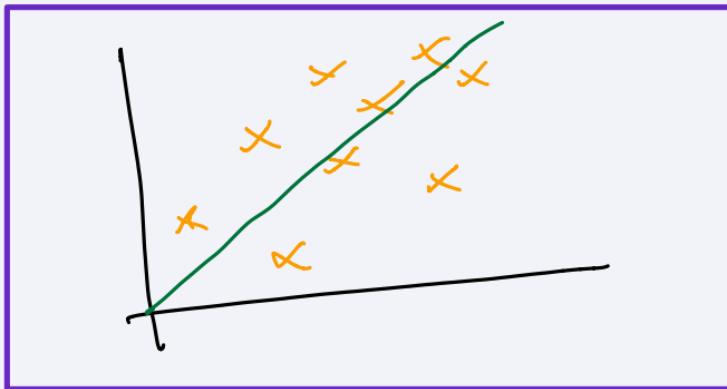
 x_1
 x_2
 x_3

Linear models

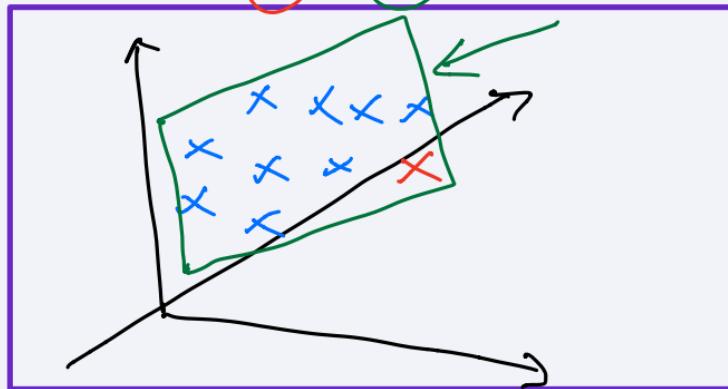
Higher dimensional data, regression problem

Task: predict the height of a person, based on their weight, gender, height of their parents, etc.

$$\hat{y} = \theta x$$



$$\hat{y} = \vec{\theta}^T \vec{x}$$



Neural networks

Neural networks

2-dimensional data, classification

$$\vec{x}_i = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \quad \hat{y}_i \begin{cases} 1 \\ -1 \end{cases}$$

Neural networks

Formula representation

Linear model

$$\hat{y} = \vec{\theta}^T \vec{x} + b$$

Perception

$$\hat{y} = \underline{\sigma}(\vec{\theta}^T \vec{x})$$

activation fn

Feed forward
NNs

$$\hat{y} = \underline{\sigma}(\underline{\underline{\theta}}_2 \sigma(\underline{\theta}_1^T \vec{x}))$$

perception

$$\hat{y} = \sigma(\underline{\theta}_3 \sigma(\underline{\theta}_2 \sigma(\underline{\theta}_1^T \vec{x}))) \leftarrow$$

Neural networks

Matrix representation

1 hidden layer FFNN

$$\hat{y} = \sigma(\theta_2 \sigma(\theta_1 \vec{x}))$$

$\sigma(\theta_1) =$

$\vec{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$

$\vec{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} h_1 \\ h_2 \\ h_3 \end{bmatrix}$

$\sigma(\theta_1 x_1 + \theta_2 x_2) = h_1$

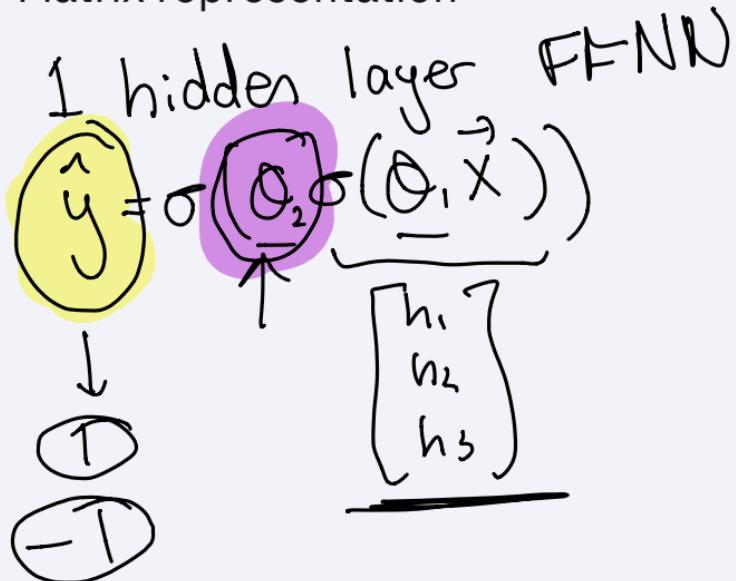
$\sigma(\theta_3 x_1 + \theta_4 x_2) = h_2$

$\sigma(\theta_5 x_1 + \theta_6 x_2) = h_3$

$\theta_1 = \begin{bmatrix} \theta_1 & \theta_2 \\ \theta_3 & \theta_4 \\ \theta_5 & \theta_6 \end{bmatrix}$

Neural networks

Matrix representation



$$\hat{y} = \sigma(\theta_2 \sigma(\theta_1 \vec{x}))$$

$1 \times 3 \quad 3 \times 1$

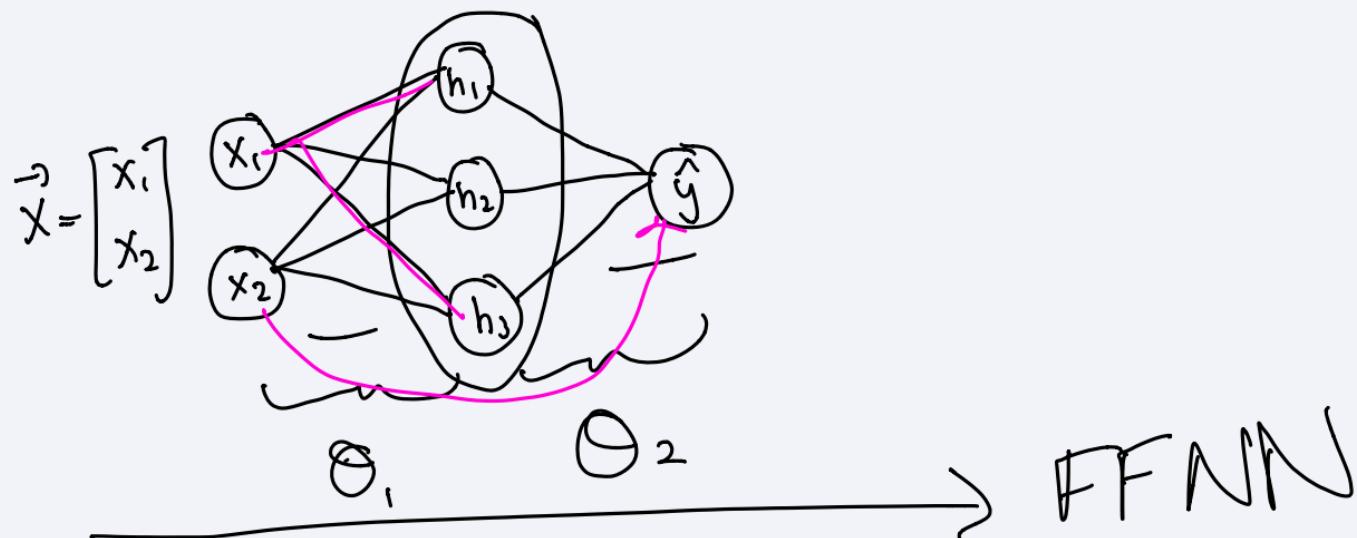
$$\theta_2 = [\theta_1' \quad \theta_2' \quad \theta_3']$$

← neurons

$$\hat{y} = \sigma(\theta_1' h_1 + \theta_2' h_2 + \theta_3' h_3)$$

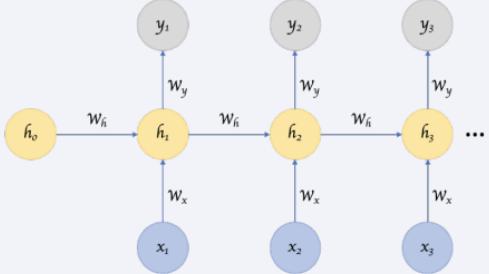
Neural networks

Graphical representation

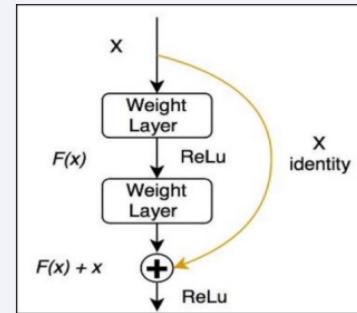


Neural networks

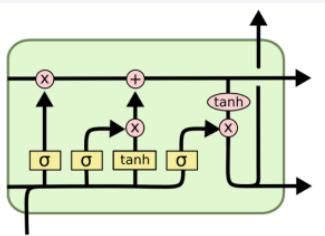
The success of neural networks



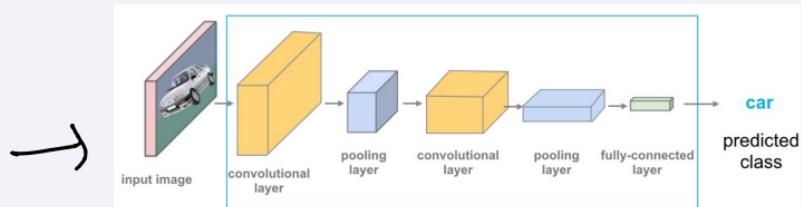
Mikolov et al. "Extensions of recurrent neural network language model." 2011 IEEE international conference on acoustics, speech and signal processing (ICASSP). IEEE, 2011.



Hao, et al. "Visualizing the loss landscape of neural nets." Advances in Neural Information Processing Systems. 2018.



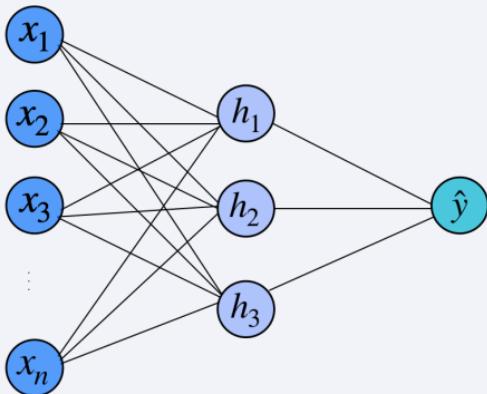
Sak et. al. "Long short-term memory recurrent neural network architectures for large scale acoustic modeling." (2014).



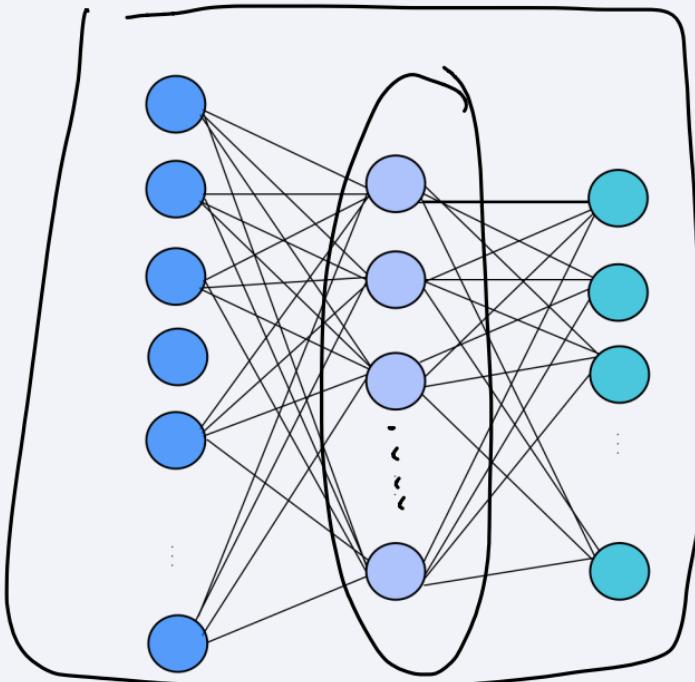
Krizhevsky et al. "Imagenet classification with deep convolutional neural networks." Communications of the ACM 60.6 (2017): 84-90.

Neural networks

The success of neural networks



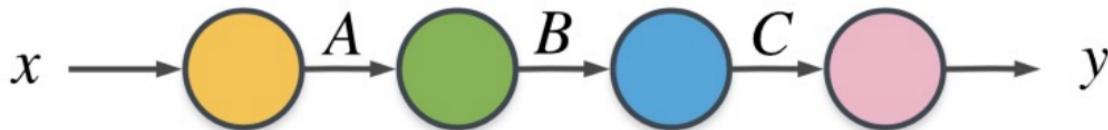
Hornik et al. "Multilayer feedforward networks are universal approximators." *Neural networks* 2.5 (1989): 359-366.



Michael Nielsen - <https://www.youtube.com/watch?v=IjgkC7OLenI>

Neural networks

- 1970 & 1986: Backpropagation, Seppo Linnainmaa & Geoffrey Hinton

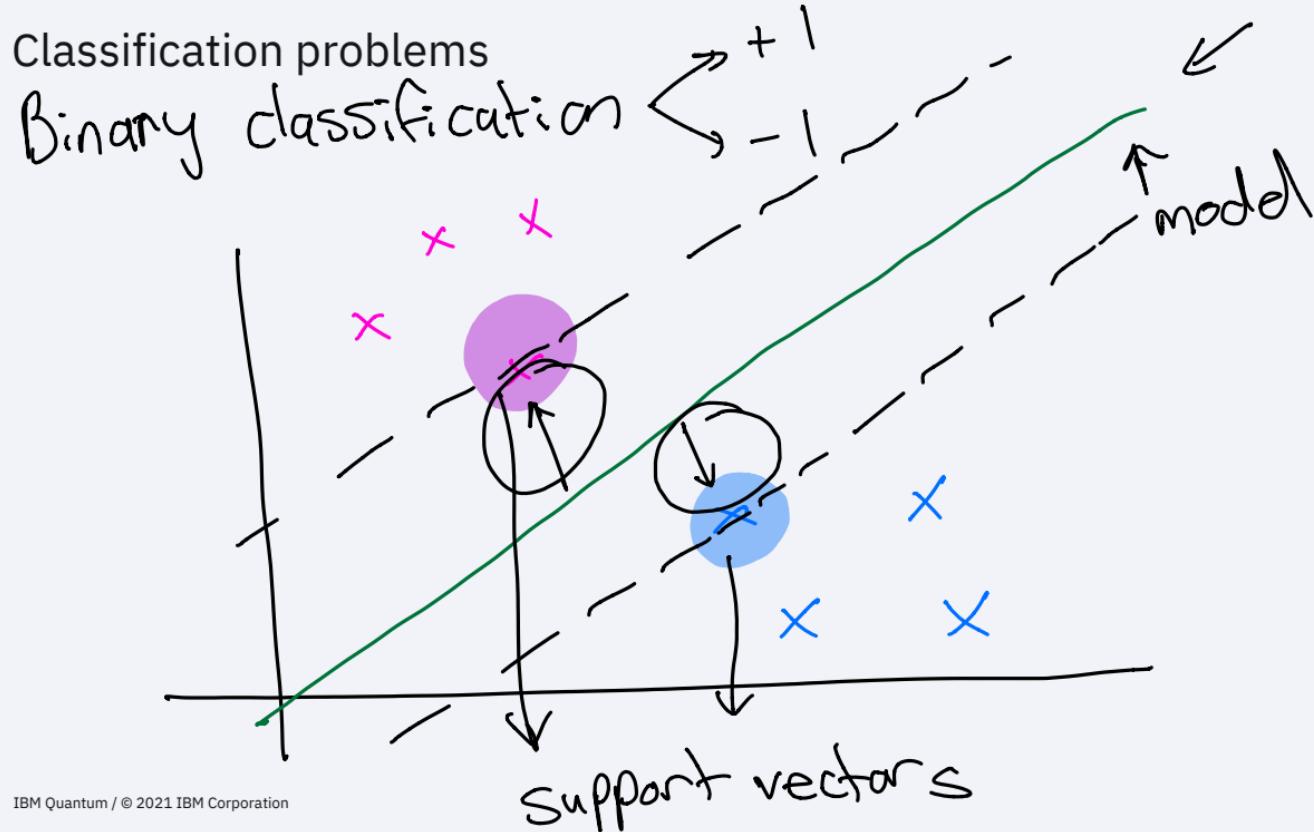


$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial C} \times \frac{\partial C}{\partial B} \times \frac{\partial B}{\partial A} \times \frac{\partial A}{\partial x}$$

Python

Support vector machines

Support vector machines



Support vector machines

Classification problems

Linear classifier

$$f(x) = \theta^T \vec{x} + b$$

$$\theta \in \mathbb{R}^d$$

d-dimensional

$$y \in [-1, 1]$$

→ quadratic opt.

- Primal formulation

Support vector machines

Classification problems

- Dual formulation

$$f(x) = \sum_{i=1}^n \alpha_i y_i (x_i^T x) + b$$

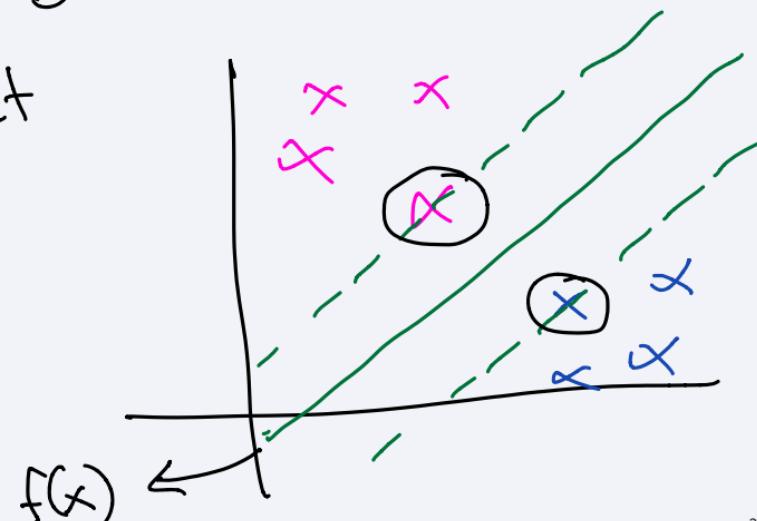
inner product
 x_i x

Why?

most $\alpha_i = 0$

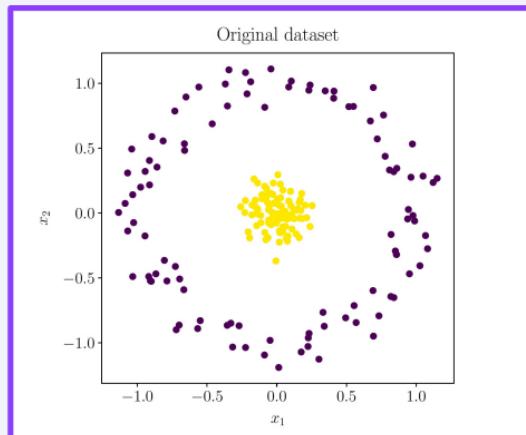
$\# \alpha_i \neq 0 \rightarrow \# \text{support vectors}$

$$f(x) = \theta^T x + b$$



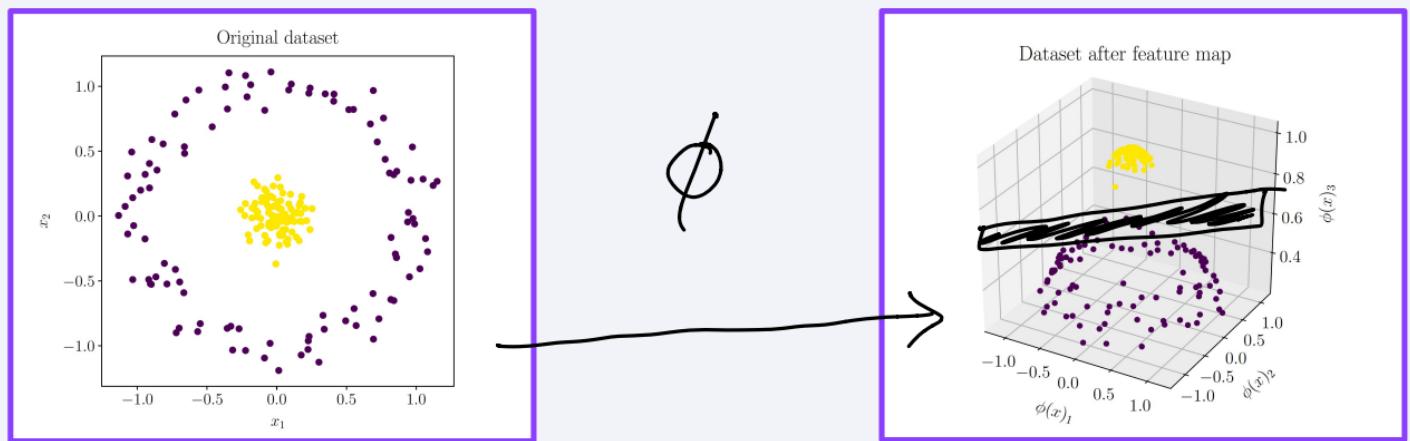
Support vector machines

Classification problems – data not linearly separable?



Support vector machines

Classification problems – data not linearly separable? Apply a **feature map**



$$\vec{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

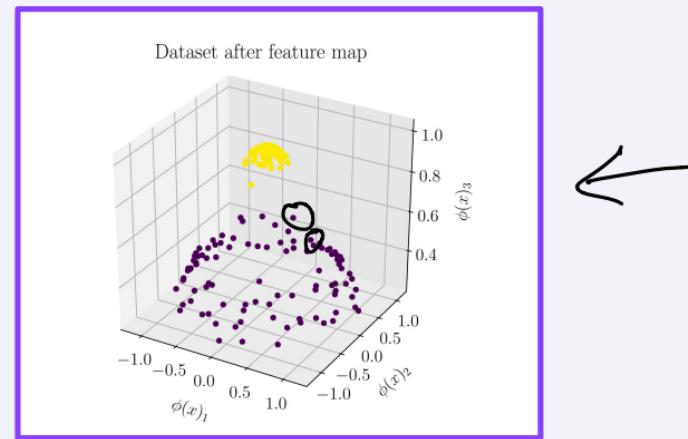
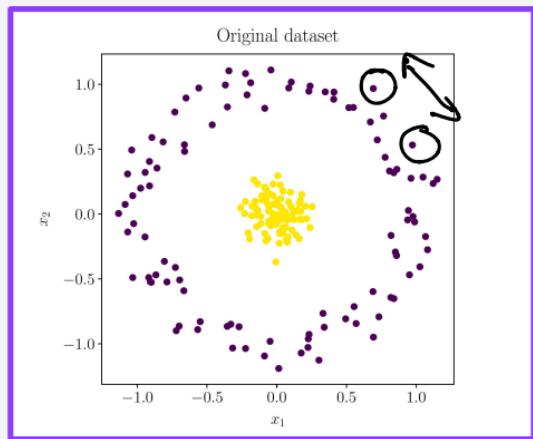
2-d

$$\phi(\vec{x}) = \begin{bmatrix} x_1 \\ x_2 \\ x_1 x_2 \end{bmatrix}$$

3d

Support vector machines

Classification problems – data not linearly separable? Apply a **feature map**



$$\frac{x_i^T x_j}{\text{similarity}}$$

$$\phi(\vec{x})$$

$$\underbrace{\phi(x_i)^T \phi(x_j)}_{\text{similarity}}$$

Support vector machines

Classification problems

$$f(x) = \sum_i \alpha_i y_i (x_i^T x) + b$$

$$f(x) = \theta^T x + b$$

$$f(x) = \sum_i \alpha_i y_i (\underbrace{\phi(x_i)^T \phi(x)}_{\text{inner product}}) + b$$

$$f(x) = \underbrace{\theta^T}_{\leftarrow} \underbrace{\phi(x)}_{\leftarrow} + b \quad \nwarrow \quad \nearrow \text{kernel}$$

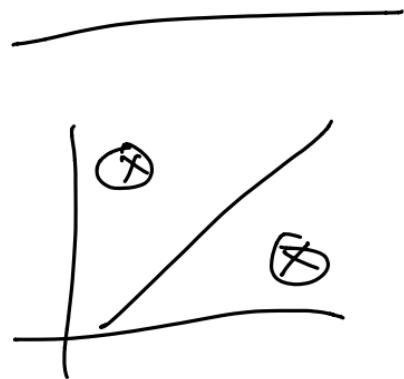
Recap

Linear models

$$\hat{y} = \theta^T x + b$$

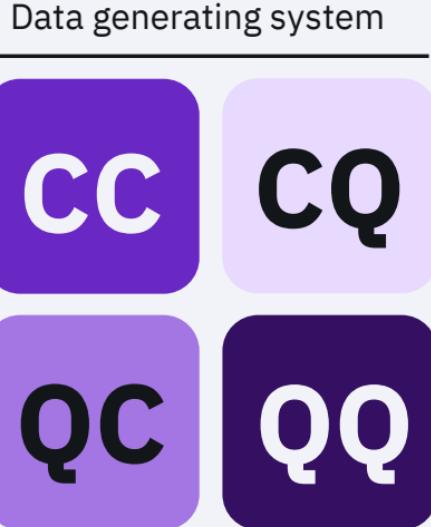

Neural networks

Support vector machines



Quantum machine learning

Data processing device



C – *classical*, Q – *quantum*

Image credit: Maria Schuld and Francesco Petruccione. *Supervised learning with quantum computers*. Vol. 17. Springer, 2018.

Quantum machine learning

Data processing device



C – *classical*, Q – *quantum*

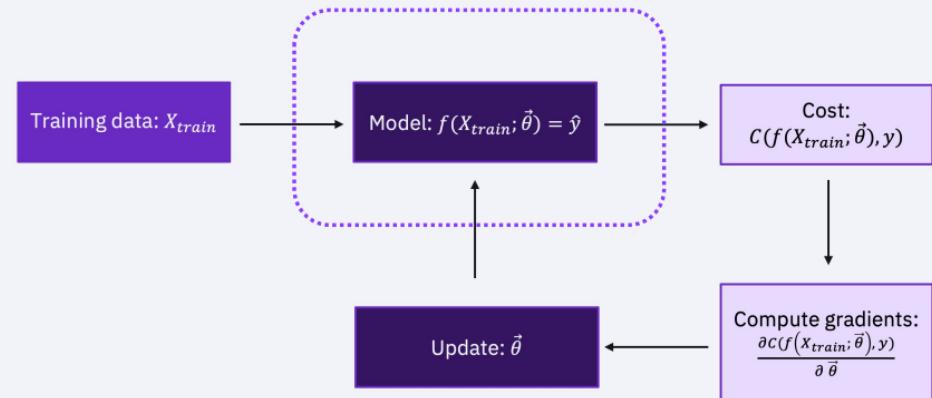


Image credit: Maria Schuld and Francesco Petruccione. *Supervised learning with quantum computers*. Vol. 17. Springer, 2018.

For the next lectures

For the next lectures, brush up on classical ML and intro to QC concepts.

In particular:

- Linear algebra basics (vectors, matrices, tensors) ←
 - Dirac notation (bra-ket notation, measurement operators) ←
 - Circuit notation (gates) ←
 - Cost function minimization (backpropagation, gradient descent) ←
 - Support vector machines (data, feature maps, kernels) ←
- 

Thank you!



Amira Abbas

@AmiraMorphism